

PURE 4-Document Cross-Model Paper Report

Generated: 2026-04-29T03:29:51.415807+00:00

Gemini-2.5-Pro-FullContext: gemini-2.5-pro (gemini)

Ollama-Qwen2.5-7B-FullContext: qwen2.5:7b-instruct (ollama)

Included Documents

Sample ID	Document ID	Source Requirements
pure_0000_cctns	0000_cctns	100
pure_0000_gamma_j	0000_gamma_j	57
pure_1999_dii	1999_dii	22
pure_2005_microcare	2005_microcare	25

Headline Findings

Higher dialogue recall: Gemini-2.5-Pro-FullContext 0.9118 vs Ollama-Qwen2.5-7B-FullContext 0.8824
Stronger direct baseline F1: Gemini-2.5-Pro-FullContext 0.7784 vs Ollama-Qwen2.5-7B-FullContext 0.9203
Stronger conversational semantic recall: Gemini-2.5-Pro-FullContext 0.6912 vs Ollama-Qwen2.5-7B-FullContext 0.3725
Stronger conversational semantic F1: Gemini-2.5-Pro-FullContext 0.7032 vs Ollama-Qwen2.5-7B-FullContext 0.4258
Stronger weighted conversational validator recall: Gemini-2.5-Pro-FullContext 0.5294 vs Ollama-Qwen2.5-7B-FullContext 0.3211

Aggregate Headline Metrics

Model	Dialogue R	Direct P	Direct R	Direct F1	Pipeline P	Pipeline R	Pipeline F1	LLM Weighted R	LLM Weighted F1
Gemini-2.5-Pro-FullContext	0.9118	0.9257	0.6716	0.7784	0.7157	0.6912	0.7032	0.5294	0.5387
Ollama-Qwen2.5-7B-FullContext	0.8824	0.9676	0.8775	0.9203	0.4967	0.3725	0.4258	0.3211	0.3669

Per-Document Dialogue Coverage

Sample ID	Reqs	Gemini-2.5-Pro-FullContext Recall	Gemini-2.5-Pro-FullContext Units	Ollama-Qwen2.5-7B-FullContext Recall	Ollama-Qwen2.5-7B-FullContext Units
pure_0000_cctns	100	0.9400	105	0.9000	128
pure_0000_gamma_j	57	0.8596	60	0.8246	88
pure_1999_dii	22	0.8636	32	0.8636	64
pure_2005_microcare	25	0.9600	38	0.9600	47

Per-Document Direct Semantic Metrics

ID	Gemini-2.5-Pro-FullContext P	Gemini-2.5-Pro-FullContext R	Gemini-2.5-Pro-FullContext F1	Ollama-Qwen2.5-7B-FullContext P	Ollama-Qwen2.5-7B-FullContext R	Ollama-Qwen2.5-7B-FullContext F1
pure_0000_cctns	0.9592	0.4700	0.6309	0.9600	0.9600	0.9600
pure_0000_gamma_j	0.9630	0.9123	0.9369	1.0000	0.7895	0.8824
pure_1999_dii	0.8947	0.7727	0.8293	1.0000	0.7727	0.8718
pure_2005_microcare	0.8077	0.8400	0.8235	0.9130	0.8400	0.8750

Per-Document Conversational Semantic Metrics

ID	Gemini-2.5-Pro-FullContext P	Gemini-2.5-Pro-FullContext R	Gemini-2.5-Pro-FullContext F1	Ollama-Qwen2.5-7B-FullContext P	Ollama-Qwen2.5-7B-FullContext R	Ollama-Qwen2.5-7B-FullContext F1
pure_0000_cctns	0.7912	0.7200	0.7539	0.5775	0.4100	0.4795
pure_0000_gamma_j	0.7292	0.6140	0.6667	0.4286	0.3684	0.3962
pure_1999_dii	0.5909	0.5909	0.5909	0.1429	0.0909	0.1111
pure_2005_microcare	0.5833	0.8400	0.6885	0.6316	0.4800	0.5455

Per-Document Conversational Weighted Validator Metrics

Sample ID	Gemini-2.5-Pro-FullContext Weighted R	Gemini-2.5-Pro-FullContext Weighted F1	Ollama-Qwen2.5-7B-FullContext Weighted R	Ollama-Qwen2.5-7B-FullContext Weighted F1
pure_0000_cctns	0.5700	0.5969	0.3500	0.4094
pure_0000_gamma_j	0.5351	0.5810	0.3333	0.3585
pure_1999_dii	0.3409	0.3409	0.2045	0.2500
pure_2005_microcare	0.5200	0.4262	0.2800	0.3182

Per-Document Conversational Diagnostics

ni-2.5-Pro-FullContext Evidence	Gemini-2.5-Pro-FullContext Props	Gemini-2.5-Pro-FullContext Grounded	Ollama-Qwen2.5-7B-FullContext Evidence	Ollama-Qwen2.5-7B-FullContext Props	Ollama-Qwen2.5-7B-FullContext Grounded
	91	91	128	71	71
	48	48	88	49	49
	22	22	64	14	14
	36	36	47	19	19

Run Paths

Gemini-2.5-Pro-FullContext: outputs/pure_full_runs/20260429T001722Z_pure_full
Ollama-Qwen2.5-7B-FullContext: outputs/pure_full_runs/20260429T002305Z_pure_full