

Benchmark run report

Run ID: 20260428T201035Z_pure_full
Date generated: 2026-04-28 23:52:04
Generation model: gemini / gemini-3.1-pro-preview
Standard validator: gemini / gemini-2.5-flash
Dialogue method: g_adaptive_semantic_coverage_v2
Question algorithm: semantic_gap_llm_v2
Clarification rounds: requested 2 | used 2

Layer 1 - Sentence-transformer matching

Condition	Precision	Recall	F1	Hallucination rate
Direct baseline	1.0000	1.0000	1.0000	0.0000
Conversational pipeline	0.9403	0.6300	0.7545	0.0597
Oracle	1.0000	1.0000	1.0000	0.0000

Layer 2 - Gemini checklist judge

Condition	Strict P	Strict R	Strict F1	Weighted P	Weighted R	Weighted F1
Direct baseline	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000
Conversational pipeline	0.5224	0.3500	0.4192	0.6940	0.4650	0.5569

Judge scoring: strict counts full=1, partial=0, none=0. Weighted counts full=1, partial=0.5, none=0.

Per-document coverage

0000_cctns: 63 of 100 requirements recovered (63.0%)



Dialogue upper bound: 90.0% of source requirements present in user turns
Evidence units: 103 | Propositions: 67 | Gap-pass additions: 4 | Grounded after rewrite: 67

Key finding: Direct baseline outperforms conversational pipeline. Dialogue stage may be lossy.