


Benchmark run report

Run ID: 20260428T205422Z_pure_full
Date generated: 2026-04-29 00:37:48
Generation model: ollama / qwen2.5:7b-instruct
Standard validator: not configured
Dialogue method: g_adaptive_semantic_coverage_v2
Question algorithm: semantic_gap_llm_v2
Clarification rounds: requested 2 | used 2

Layer 1 - Sentence-transformer matching

Condition	Precision	Recall	F1	Hallucination rate
Direct baseline	0.9600	0.9600	0.9600	0.0400
Conversational pipeline	0.6515	0.4300	0.5181	0.3485
Oracle	1.0000	1.0000	1.0000	0.0000

Per-document coverage



Dialogue upper bound: 90.0% of source requirements present in user turns
Evidence units: 128 | Propositions: 66 | Gap-pass additions: 4 | Grounded after rewrite: 66
Key finding: Direct baseline outperforms conversational pipeline. Dialogue stage may be lossy.