

Benchmark run report

Run ID: 20260429T001722Z_pure_full
Date generated: 2026-04-29 04:54:43
Generation model: gemini / gemini-2.5-pro
Standard validator: gemini / gemini-2.5-flash
Dialogue method: g_adaptive_semantic_coverage_v2
Question algorithm: semantic_gap_llm_v2
Clarification rounds: requested 2 | used 0-2

Layer 1 - Sentence-transformer matching


Condition	Precision	Recall	F1	Hallucination rate
Direct baseline	0.9257	0.6716	0.7784	0.0939
Conversational pipeline	0.7157	0.6912	0.7032	0.3263
Oracle	1.0000	1.0000	1.0000	0.0000

Layer 2 - Gemini checklist judge

Condition	Strict P	Strict R	Strict F1	Weighted P	Weighted R	Weighted F1
Direct baseline	0.8041	0.5833	0.6761	0.8649	0.6275	0.7273
Conversational pipeline	0.3909	0.3775	0.3840	0.5482	0.5294	0.5387

Judge scoring: strict counts full=1, partial=0, none=0. Weighted counts full=1, partial=0.5, none=0.

Per-document coverage



Dialogue upper bound: 91.2% of source requirements present in user turns
Evidence units: 235 | Propositions: 197 | Gap-pass additions: 7 | Grounded after rewrite: 197
Key finding: Direct baseline outperforms conversational pipeline. Dialogue stage may be lossy.