

Benchmark run report

Run ID: 20260429T002305Z_pure_full
Date generated: 2026-04-29 05:35:58
Generation model: ollama / qwen2.5:7b-instruct
Standard validator: gemini / gemini-2.5-flash
Dialogue method: g_adaptive_semantic_coverage_v2
Question algorithm: semantic_gap_llm_v2
Clarification rounds: requested 2 | used 0-2

Layer 1 - Sentence-transformer matching

Condition	Precision	Recall	F1	Hallucination rate
Direct baseline	0.9676	0.8775	0.9203	0.0317
Conversational pipeline	0.4967	0.3725	0.4258	0.5549
Oracle	1.0000	1.0000	1.0000	0.0000

Layer 2 - Gemini checklist judge

Condition	Strict P	Strict R	Strict F1	Weighted P	Weighted R	Weighted F1
Direct baseline	0.8595	0.7794	0.8175	0.9243	0.8382	0.8792
Conversational pipeline	0.1961	0.1471	0.1681	0.4281	0.3211	0.3669

Judge scoring: strict counts full=1, partial=0, none=0. Weighted counts full=1, partial=0.5, none=0.

Per-document coverage

0000_cctns: 41 of 100 requirements recovered (41.0%)



0000_gamma_j: 21 of 57 requirements recovered (36.8%)



1999_dii: 2 of 22 requirements recovered (9.1%)



2005_microcare: 12 of 25 requirements recovered (48.0%)



Dialogue upper bound: 88.2% of source requirements present in user turns

Evidence units: 327 | Propositions: 153 | Gap-pass additions: 0 | Grounded after rewrite: 153

Key finding: Direct baseline outperforms conversational pipeline. Dialogue stage may be lossy.